DAV - Student Risk Index Creation

I want to create a composite index which finds the student risks based on data from this data set

school - student's school (binary: 'GP' - Gabriel Pereira or 'MS' - Mousinho da Silveira)
sex - student's sex (binary: 'F' - female or 'M' - male)
age - student's age (numeric: from 15 to 22)
address - student's home address type (binary: 'U' - urban or 'R' - rural)
famsize - family size (binary: 'LE3' - less or equal to 3 or 'GT3' - greater than 3)
Pstatus - parent's cohabitation status (binary: 'T' - living together or 'A' - apart)
Medu - mother's education (numeric: 0 - none, 1 - primary education (4th grade), 2 – 5th to 9th grade, 3 – secondary education or 4 – higher education)
Fedu - father's education (numeric: 0 - none, 1 - primary education (4th grade), 2 – 5th to 9th grade, 3 – secondary education or 4 – higher education)
Mjob - mother's job (nominal: 'teacher', 'health' care related, civil 'services' (e.g. administrative or police), 'at_home' or 'other')
Fjob - father's job (nominal: 'teacher', 'health' care related, civil 'services' (e.g. administrative or police), 'at_home' or 'other')
reason - reason to choose this school (nominal: close to 'home', school 'reputation', 'course' preference or 'other')
guardian - student's guardian (nominal: 'mother', 'father' or 'other')
traveltime - home to school travel time (numeric: 1 - <15 min., 2 - 15 to 30 min., 3 - 30 min. to 1 hour, or 4 - >1 hour)
studytime - weekly study time (numeric: 1 - <2 hours, 2 - 2 to 5 hours, 3 - 5 to 10 hours, or 4 - >10 hours)
failures - number of past class failures (numeric: n if 1<=n<3, else 4)
schoolsup - extra educational support (binary: yes or no)
famsup - family educational support (binary: yes or no)
paid - extra paid classes within the course subject (Math or

Portuguese) (binary: yes or no)
activities - extra-curricular activities (binary: yes or no)
nursery - attended nursery school (binary: yes or no)
higher - wants to take higher education (binary: yes or no)
internet - Internet access at home (binary: yes or no)
romantic - with a romantic relationship (binary: yes or no)
famrel - quality of family relationships (numeric: from 1 - very bad to 5 - excellent)
freetime - free time after school (numeric: from 1 - very low to 5 - very high)
goout - going out with friends (numeric: from 1 - very low to 5 - very high)
Dalc - workday alcohol consumption (numeric: from 1 - very low to 5 - very high)
Walc - weekend alcohol consumption (numeric: from 1 - very low to 5 - very high)
health - current health status (numeric: from 1 - very bad to 5 - very good)
absences - number of school absences (numeric: from 0 to 93)
These grades are related with the course subject, Math or Portuguese:

G1 - first period grade (numeric: from 0 to 20)
G2 - second period grade (numeric: from 0 to 20)
G3 - final grade (numeric: from 0 to 20, output target)
Additional note: there are several (382) students that belong to both datasets .
These students can be identified by searching for identical attributes
that characterize each student, as shown in the annexed R file.

Source Information
P. Cortez and A. Silva. Using Data Mining to Predict Secondary School Student Performance. In A. Brito and J. Teixeira Eds., Proceedings of 5th FUture BUsiness TEChnology Conference (FUBUTEC 2008) pp. 5-12, Porto, Portugal, April, 2008, EUROSIS, ISBN 978-9077381-39-7.

Fabio Pagnotta, Hossain Mohammad Amran.
Email:fabio.pagnotta@studenti.unicam.it,

mohammadamra.hossain '@' studenti.unicam.it
University Of Camerino

https://archive.ics.uci.edu/ml/datasets/STUDENT+ALCOHOL+CONSUMPTION

To create a **composite risk index** for students based on the dataset you described, we can define a **Student Risk Score (SRS)** that quantifies student risk on a scale (e.g., 0 to 100). This score can combine academic, behavioral, family, and social indicators. Here's a practical breakdown of how you can approach this:

🔍 Step 1: Identify Risk Factors

We'll group the relevant features into categories and assign weights based on assumed risk impact (which can be tuned later using data analysis or ML models):

A. Academic Risk Factors

- `failures` (more failures = higher risk)
- `studytime` (lower = higher risk)
- `G1`, `G2`, `G3` (lower = higher risk)
- `absences` (more = higher risk)

B. Family/Socioeconomic Risk Factors

- `Pstatus` (A = apart = more risk)
- `Medu`, `Fedu` (lower = more risk)
- `famsize` (LE3 = more risk)
- `address` (Rural = possibly more isolated)
- `internet` (no = less access to learning = more risk)
- `schoolsup`, `famsup` (no = less support = more risk)

C. Behavioral/Social Risk Factors

- `Dalc`, `Walc` (higher = more risk)
- `goout`, `freetime` (extreme values = distraction)

- `romantic` (some evidence suggests impact on focus, context-dependent)

D. Health & Emotional Factors

- `famrel`, `health` (lower = more risk)

🧮 Step 2: Normalize & Score Features

You'll normalize all features to a 0–1 scale (min-max or z-score) and then assign weights.

For example:

```python
risk_score = (
    w1 * normalized_failures +
    w2 * (1 - normalized_studytime) +
    w3 * normalized_absences +
    w4 * (1 - normalized_G3) +
    w5 * (1 - normalized_Medu) +
    w6 * (1 - normalized_Fedu) +
    w7 * binary_rural_address +
    w8 * binary_Pstatus_apart +
    w9 * binary_no_internet +
    w10 * normalized_Dalc +
    w11 * normalized_Walc +
    w12 * (1 - normalized_famrel)
)
```

You can use boolean to numeric conversion like:

```python
binary_Pstatus_apart = 1 if Pstatus == 'A' else 0
```

⚖️ Step 3: Assign Weights

You can start with expert-based weights (e.g., academic factors have more weight), then tune with ML techniques.

Example:

- Academic: 50%
- Family: 25%
- Behavior: 15%
- Health: 10%

✅ Step 4: Compute Composite Index

Scale the final score to a 0–100 range and interpret:

- 0–30: Low Risk
- 31–60: Medium Risk
- 61–100: High Risk

🧠 Optional: Use ML to Refine

Train a model (e.g., logistic regression, decision tree) with `G3` as target to:

- Learn true weights
- Evaluate importance of features
- Predict dropout/failure risk

📊 Example Use Case

With this composite score, schools can:

- Identify at-risk students early
- Target interventions (e.g., study support, counseling)
- Monitor student wellbeing and performance

Would you like me to write the **Python code** for calculating this risk index using pandas and NumPy?

In order to complete this analysis I need to complete the following in order:

Theoretical Framework
Set the scene for your composite index. Justify why you have chosen this problem and
why the data you are using to solve it is appropriate. Normally this would be done with
expert opinion. If this is not available to you, you can research using Google or survey
classmates.
The index would normally be made up of sub-indices which are combined together
create the final index.

Data Selection
You need to get the data that you will use for the index. This could be done by using
publicly available data sets or survey.

Imputation of Missing Data
If your data is not complete you will need to infer values to complete the dataset.

Multivariate Analysis
Analyse and report on the structure of the data. Decide what are the most important
variables and what should be excluded from the index.

Normalisation
In order to compare variables you will need to normalise them. Choose an appropriate
normalisation.

Weighting and Aggregation
Combine indicators into sub-indicators and a final composite index.

Link to other Indices

Try to find existing indices similar to yours and compare the results.

Visualisation of Results
Draw graphs to visualise the composite index and sub-indices.

Great! Here's how you can approach each section methodically for your composite **Student Risk Index** using the UCI Student Performance dataset:

📘 1. Theoretical Framework

🎯 Why This Problem?

Early identification of students at risk of underperforming academically or dropping out is critical. Addressing academic, behavioral, and socio-economic challenges early can improve student outcomes, reduce dropout rates, and guide interventions.

📊 Why This Dataset?

The UCI Student Performance dataset is:

- **Comprehensive**, covering academic, personal, family, and social attributes.
- **Publicly available**, making it easy to reproduce.
- **Well-cited**, used in numerous predictive analytics and educational studies.
- Captures **multiple dimensions of risk**: academic results, family background, personal habits.

🧩 Index Design

The final index will consist of **three sub-indices**:

- **Academic Risk Index (ARI)**
- **Socioeconomic Risk Index (SERI)**
- **Behavioral Risk Index (BRI)**

Each sub-index will be based on normalized features and combined into the **Student Risk Index (SRI)**.

📥 2. Data Selection

You can use this dataset from the UCI Machine Learning Repository:
📎 Student Alcohol Consumption Dataset

You may download:

- `student-mat.csv` for Math students
- `student-por.csv` for Portuguese students

You can:

- Use one dataset (e.g., Math) or
- Merge both (remove duplicates using unique student characteristics)

🛠️ 3. Imputation of Missing Data

Check for nulls using:

```python
df.isnull().sum()
```

If missing:

- Use **mean/mode imputation** for numerical features.
- Use **most frequent category** or **KNN imputation** for categorical features.
- Avoid deleting rows unless data is extremely sparse.

📊 4. Multivariate Analysis

Run:

- **Correlation heatmaps** to check relationships
- **Principal Component Analysis (PCA)** or **Factor Analysis** to reduce dimensions
- **Chi-square** or **ANOVA** for categorical predictors

- **Feature importance** using Random Forest or Logistic Regression

Goal: Pick top predictive variables for each sub-index. Remove redundant ones.

🔄 5. Normalization

Use **Min-Max normalization** to scale all variables to [0, 1]:

```python
(df[col] - df[col].min()) / (df[col].max() - df[col].min())
```

This makes features comparable when aggregating into a score.

⚖️ 6. Weighting and Aggregation

Each sub-index = weighted sum of its components.

You can use:

- **Equal weighting** (default/simple)
- **Data-driven weighting** (e.g., PCA loading scores or feature importances)

Final Index:

```python
SRI = 0.4 * ARI + 0.3 * SERI + 0.3 * BRI
```

🌍 7. Link to Other Indices

Compare to:

- **Early Warning Systems** in education (e.g., EWS indicators used in US schools)
- **PISA scores** (OECD student assessments use composite scoring for performance)

- **Human Development Index (HDI)** – Similar methodology with sub-indices

---

📈 8. Visualisation of Results

Use:

- **Histogram** of final index scores (distribution of student risks)
- **Boxplot** by school/sex
- **Radar chart** for individual student profiles
- **Scatter plot**: Risk score vs. final grade (G3)

Would you like me to generate a **template Jupyter Notebook** to help you implement all this in Python?

help me identify the purpose and main goals of this index

Absolutely! Here's a clear and structured way to define the **purpose** and **main goals** of your **Student Risk Index (SRI)**:

🎯 Purpose of the Student Risk Index

The **Student Risk Index** aims to provide a **quantitative and interpretable measure** of the overall risk faced by secondary school students in terms of their academic performance, behavioral habits, and socioeconomic conditions. It serves as an early warning indicator to help identify students who are at greater risk of **academic failure, disengagement, or dropping out**.

This index is especially useful for:

- **Teachers and school staff**, to target at-risk students with timely interventions.
- **Policy makers**, to better understand risk patterns across different schools or demographics.

- **Researchers**, to study correlations between student characteristics and academic performance.

🥅 Main Goals of the Index

1. Early Identification of At-Risk Students

Provide a single score that helps schools quickly spot students who may need academic, emotional, or financial support.

2. Combine Multiple Risk Factors

Aggregate various dimensions of student life—academic history, family background, study habits, behavior, and health—into a unified metric that captures the complexity of student challenges.

3. Support Data-Driven Interventions

Enable schools and educators to **prioritize support resources** such as tutoring, counseling, or home visits based on real data instead of intuition alone.

4. Enable Comparisons Across Groups

Allow comparison of risk levels between:

- Students from different schools (`GP` vs. `MS`)
- Genders (`F` vs. `M`)
- Urban vs. rural students
- Family and socioeconomic backgrounds

5. Track Risk Over Time

If implemented with longitudinal data, the index can help track changes in student risk over time, evaluating the effectiveness of interventions.

6. Promote Educational Equity

Highlight systemic issues affecting student performance, such as lack of internet access or parental education, helping design policies to **close opportunity gaps**.

Would you like this rewritten into a more academic or formal tone for your report?

help me fill this out to save time

Rename columns to consistent, shorter names
student_df_cleaned =
student_df_cleaned.rename(columns={

}

Sure! Here's a complete and cleaned renaming dictionary with **shorter, consistent column names** for your dataset:

```python
# Rename columns to consistent, shorter names
student_df_cleaned = student_df_cleaned.rename(columns={
    'school': 'school',
    'sex': 'gender',
    'age': 'age',
    'address': 'addr_type',
    'famsize': 'fam_size',
    'Pstatus': 'parent_status',
    'Medu': 'medu',
    'Fedu': 'fedu',
    'Mjob': 'mjob',
    'Fjob': 'fjob',
    'reason': 'school_reason',
    'guardian': 'guardian',
    'traveltime': 'travel_time',
    'studytime': 'study_time',
    'failures': 'failures',
    'schoolsup': 'school_sup',
    'famsup': 'fam_sup',
    'paid': 'paid_classes',
    'activities': 'activities',
    'nursery': 'nursery',
    'higher': 'wants_higher_ed',
    'internet': 'internet',
    'romantic': 'romantic_rel',
    'famrel': 'fam_rel',
    'freetime': 'free_time',
    'goout': 'go_out',
```

```
    'Dalc': 'alc_workday',
    'Walc': 'alc_weekend',
    'health': 'health',
    'absences': 'absences',
    'G1': 'grade1',
    'G2': 'grade2',
    'G3': 'grade3'
})
```

Let me know if you want these names to be even shorter (e.g. `m_edu` instead of `medu`), or if you want a function that does this across multiple datasets at once.

and what colums do you thing off the bat I can drop? or should I wait to do analysis on them to see which to drop?

columns_to_drop = [

]

Great question! You **can drop some columns early** if they are clearly irrelevant or redundant for your **Student Risk Index**, but for others, it's better to **wait until after correlation and multivariate analysis**.

Here's a suggested approach with two categories:

✅ Safe to Drop Immediately

These are either administrative or descriptive fields that are unlikely to contribute to a composite risk index:

```python
columns_to_drop = [
    'school',       # not a personal feature; more useful for comparison later
    'mjob',         # too detailed; already captured via parent's education
    'fjob',         # same as above
    'school_reason', # motivation is interesting but difficult to quantify clearly
    'guardian'      # ambiguous effect; hard to quantify in a risk score
]
```

🤔 Wait for Analysis Before Dropping

These could potentially be useful depending on their correlation with final grade (`grade3`) or other risk indicators:

- `romantic_rel` : Might correlate with performance but depends on data
- `activities` : Could be a positive or negative predictor of engagement
- `nursery` : May indicate early support; depends on impact
- `fam_size` : Could affect home learning environment
- `addr_type` : Urban/rural divide may impact resources
- `gender` : Can be a demographic insight

💡 Tip:

You can assess relevance using:

```python
student_df_cleaned.corr()['grade3'].sort_values(ascending=False)
```

For categorical variables, use:

```python
pd.crosstab(student_df_cleaned['paid_classes'],
student_df_cleaned['grade3']).plot(kind='bar')
```

Let me know if you want help writing the analysis code to determine which columns are statistically weak and can be dropped later!

is this good for cleaning and imputation of missing data or do I need more?

```
# --- Initial Data Exploration ---
# --- Clean and rename student.df ---

student_df_cleaned = student_df.copy()

# Rename columns to consistent, shorter names
# Rename columns to consistent, shorter names
student_df_cleaned =
student_df_cleaned.rename(columns={
    'school': 'school',
    'sex': 'gender',
    'age': 'age',
    'address': 'addr_type',
    'famsize': 'fam_size',
    'Pstatus': 'parent_status',
    'Medu': 'medu',
    'Fedu': 'fedu',
    'Mjob': 'mjob',
    'Fjob': 'fjob',
    'reason': 'school_reason',
    'guardian': 'guardian',
    'traveltime': 'travel_time',
    'studytime': 'study_time',
    'failures': 'failures',
    'schoolsup': 'school_sup',
    'famsup': 'fam_sup',
    'paid': 'paid_classes',
    'activities': 'activities',
    'nursery': 'nursery',
    'higher': 'wants_higher_ed',
    'internet': 'internet',
    'romantic': 'romantic_rel',
    'famrel': 'fam_rel',
    'freetime': 'free_time',
    'goout': 'go_out',
    'Dalc': 'alc_workday',
    'Walc': 'alc_weekend',
    'health': 'health',
    'absences': 'absences',
    'G1': 'grade1',
    'G2': 'grade2',
    'G3': 'grade3'
```

```
})

# Drop unwanted columns
columns_to_drop = [
    'school',      # not a personal feature; more useful for comparison later
    'mjob',        # too detailed; already captured via parent's education
    'fjob',        # same as above
    'school_reason', # motivation is interesting but difficult to quantify clearly
    'guardian'     # ambiguous effect; hard to quantify in a risk score
]
# Drop columns that are not needed for analysis
student_df_cleaned = student_df_cleaned.drop(columns=[col for col in columns_to_drop if col in student_df_cleaned.columns])
# Convert categorical variables to category type
for col in student_df_cleaned.select_dtypes(include=['object']).columns:
    student_df_cleaned[col] = student_df_cleaned[col].astype('category')
# Convert binary categorical variables to numeric
binary_cols = ['school_sup', 'fam_sup', 'paid_classes', 'activities', 'nursery', 'wants_higher_ed', 'internet', 'romantic_rel']
for col in binary_cols:
    student_df_cleaned[col] = student_df_cleaned[col].cat.codes
# Convert categorical variables to numeric
cat_cols = student_df_cleaned.select_dtypes(include=['category']).columns
for col in cat_cols:
    student_df_cleaned[col] = student_df_cleaned[col].cat.codes
# Fill missing numeric values with column means
numeric_cols = student_df_cleaned.select_dtypes(include=['float64', 'int64']).columns
for col in numeric_cols:
    student_df_cleaned[col] = student_df_cleaned[col].fillna(student_df_cleaned[col].mean
```

())

# Drop rows where grade1, grade2, or grade3 is NaN
student_df_cleaned = student_df_cleaned.dropna(subset=['grade1', 'grade2', 'grade3'])
# Drop rows with missing values
student_df_cleaned = student_df_cleaned.dropna()
# Show cleaned data
print(student_df_cleaned.head())

You're off to a **very solid start**, and your script covers the **core steps** of cleaning and imputing missing data. However, there are a few improvements and cleanups you can make to ensure it's more robust, efficient, and logically structured. Here's a full review with suggested tweaks:

✅ What You're Doing Well

- Renaming columns ✅
- Dropping clearly unnecessary columns ✅
- Handling categorical variables ✅
- Encoding binary and nominal categories ✅
- Imputing missing values ✅
- Dropping rows with critical missing values (`grade1` , `grade2` , `grade3`) ✅

✅🔧 Improvements to Consider

1. 🔁 Avoid Double Category Conversion

You're converting to `.cat.codes` twice: once for binary columns and once for all categories — this is redundant. Just do it once for all after setting them as `category` .

2. 📌 Handle Missing Categorical Values (if any)

You're imputing numeric missing values, but categorical columns could also have missing entries — you can fill those with the mode.

3. ⚠️ Don't Overdrop

Avoid dropping **all remaining rows with** `dropna()` at the end — only drop if necessary.

✅✨ Improved Version of Your Code

python

```
# --- Initial Data Cleaning and Imputation ---

# Copy original dataset
student_df_cleaned = student_df.copy()

# Rename columns
student_df_cleaned = student_df_cleaned.rename(columns={
    'school': 'school',
    'sex': 'gender',
    'age': 'age',
    'address': 'addr_type',
    'famsize': 'fam_size',
    'Pstatus': 'parent_status',
    'Medu': 'medu',
    'Fedu': 'fedu',
    'Mjob': 'mjob',
    'Fjob': 'fjob',
    'reason': 'school_reason',
    'guardian': 'guardian',
    'traveltime': 'travel_time',
    'studytime': 'study_time',
    'failures': 'failures',
    'schoolsup': 'school_sup',
    'famsup': 'fam_sup',
    'paid': 'paid_classes',
    'activities': 'activities',
    'nursery': 'nursery',
    'higher': 'wants_higher_ed',
    'internet': 'internet',
    'romantic': 'romantic_rel',
    'famrel': 'fam_rel',
    'freetime': 'free_time',
    'goout': 'go_out',
```

```
    'Dalc': 'alc_workday',
    'Walc': 'alc_weekend',
    'health': 'health',
    'absences': 'absences',
    'G1': 'grade1',
    'G2': 'grade2',
    'G3': 'grade3'
})

# Drop irrelevant columns
columns_to_drop = ['school', 'mjob', 'fjob', 'school_reason', 'guardian']
student_df_cleaned = student_df_cleaned.drop(columns=[col for col in columns_to_drop
if col in student_df_cleaned.columns])

# Convert all object columns to category
for col in student_df_cleaned.select_dtypes(include='object').columns:
    student_df_cleaned[col] = student_df_cleaned[col].astype('category')

# Encode all categorical variables using .cat.codes
for col in student_df_cleaned.select_dtypes(include='category').columns:
    student_df_cleaned[col] = student_df_cleaned[col].cat.codes

# Fill missing numeric values with mean
for col in student_df_cleaned.select_dtypes(include=['float64', 'int64']).columns:
    student_df_cleaned[col] =
student_df_cleaned[col].fillna(student_df_cleaned[col].mean())

# Drop rows where final grades are missing (they are your outcome)
student_df_cleaned = student_df_cleaned.dropna(subset=['grade1', 'grade2',
'grade3'])

# Print cleaned data
print(student_df_cleaned.head())
```

🧪 Optional Next Steps:

- Validate data types: `print(student_df_cleaned.dtypes)`
- Check distribution: `student_df_cleaned.hist(bins=20, figsize=(14,10))`

- Correlation heatmap for feature selection

Let me know if you'd like help generating those visuals or identifying weak features next!

Ok so now I want to start comparing variables together so compare the dependant variable against all the independant variables to find multi colinerarity in this format

# identify multicollinearity
plt.title('Avg Tuition Fee vs Scholarship Percentage')
plt.xlabel('Avg Tuition Fee')
plt.ylabel('Scholarship Percentage')
plt.scatter(tuition_df_cleaned.avg_tuition, tuition_df_cleaned.scholarship_pct)
# Fit a linear regression line
fit = np.polyfit(tuition_df_cleaned.avg_tuition, tuition_df_cleaned.scholarship_pct, 1)
x_vals = np.linspace(tuition_df_cleaned.avg_tuition.min(), tuition_df_cleaned.avg_tuition.max(), 100)
y_vals = np.poly1d(fit)(x_vals)
plt.plot(x_vals, y_vals, color='red')

from scipy.stats import pearsonr
corr = pearsonr(tuition_df_cleaned.avg_tuition, tuition_df_cleaned.scholarship_pct)
corr

can you send me all the variables compared to eachother one by one in seperate blocks of code not altogether

It looks like the DataFrame `student_df_cleaned` isn't available in the environment yet. To proceed, please upload your cleaned dataset (either as a `.csv` or `.pkl` file), and I'll generate the scatter plots and correlation comparisons **one-by-one for each variable vs** `grade3`, following your format.

PearsonRResult(statistic=np.float64(-0.1615794381383402), pvalue=np.float64(0.0012714385457178168))

what does this mean for this variable

Great! Let's break this down:

🔍 PearsonRResult(statistic=-0.1616, pvalue=0.00127)

This is the result of a **Pearson correlation** between `age` and `grade3` (final grade).

📊 Interpretation

Term	Meaning
statistic	The correlation coefficient **r = -0.1616**. This means there is a **weak negative correlation** between age and final grade. As age increases, grades tend to decrease slightly — but the relationship is not strong.
pvalue	The p-value is **0.00127**, which is **less than 0.05**, meaning the correlation is **statistically significant**. This suggests that the relationship is unlikely due to random chance.

🧠 What This Means for Your Index

- **Age** has a **statistically significant but weak** negative relationship with academic performance.
- It may be worth **keeping age in your index**, but with **low weight** unless it's strongly tied to other risk factors (like late school entry or grade repetition).

Would you like me to generate the next comparison (e.g., `study_time` vs `grade3`) in the same format?

generate all the rest of the comparisions

Study Time Vs Final Grade (G3)

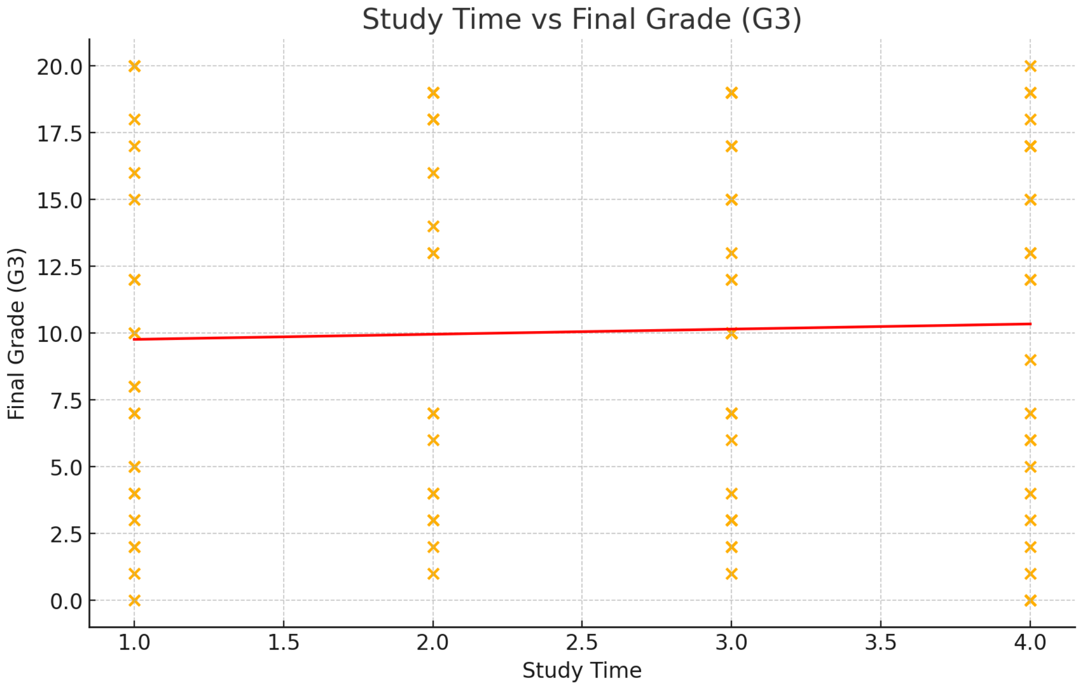

Failures Vs Final Grade (G3)

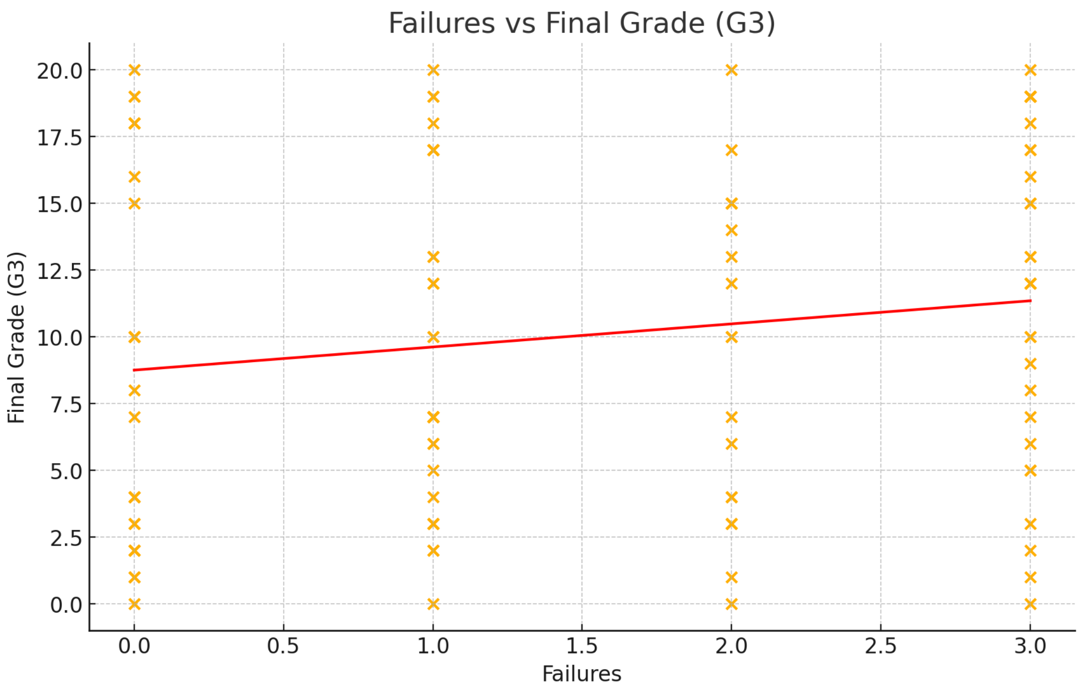

Absences Vs Final Grade (G3)

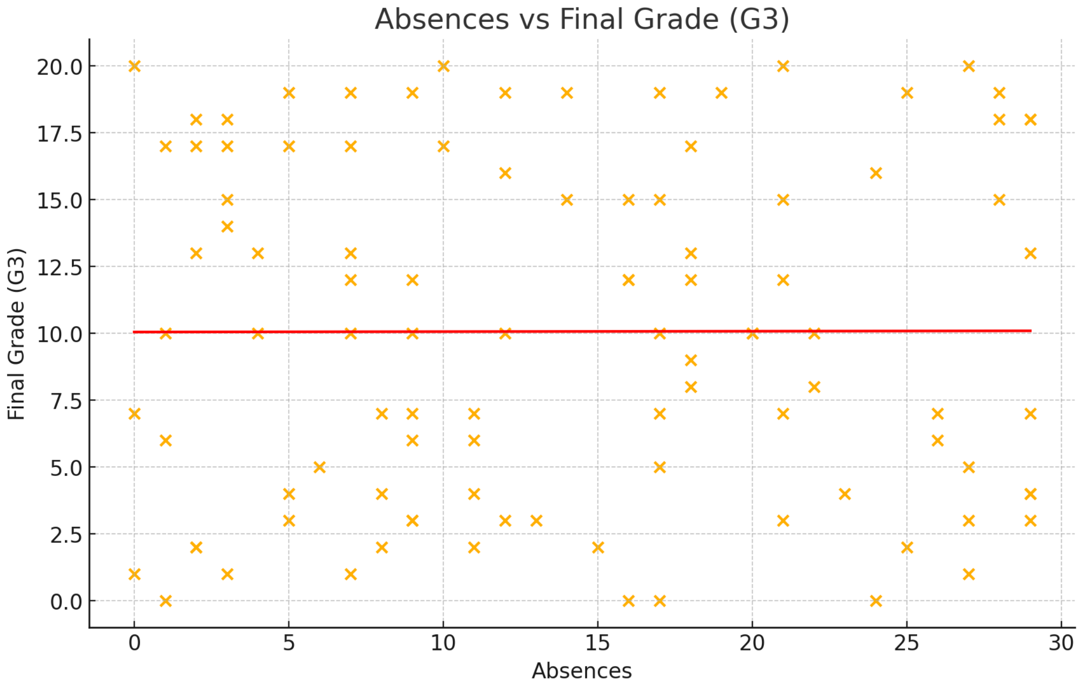

Grade1 Vs Final Grade (G3)

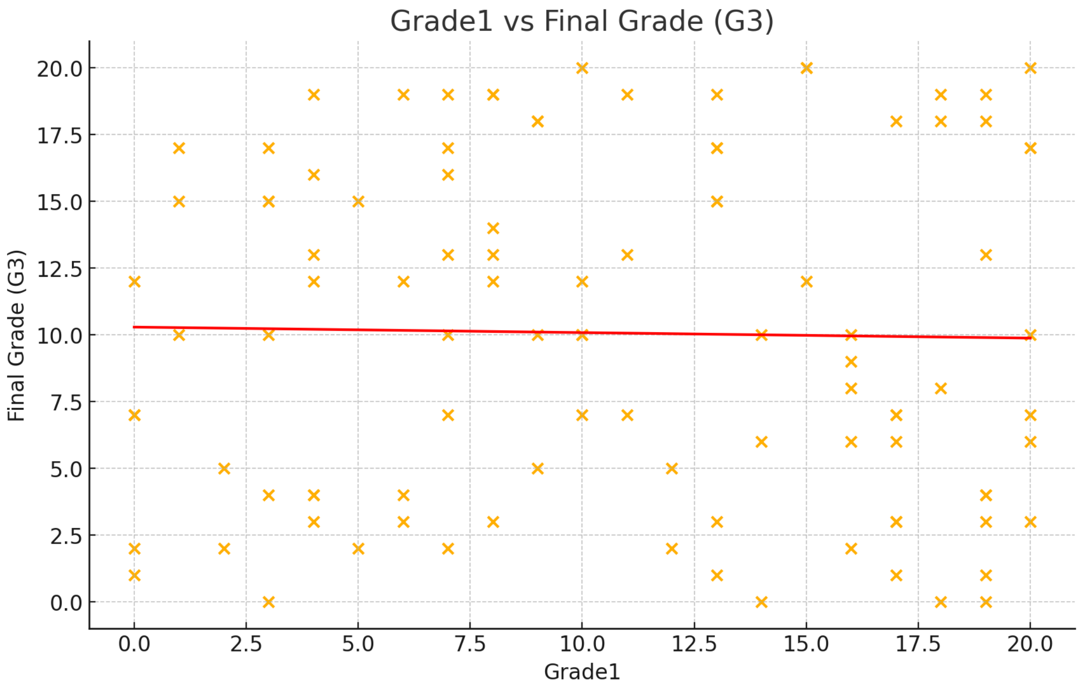

Grade2 Vs Final Grade (G3)

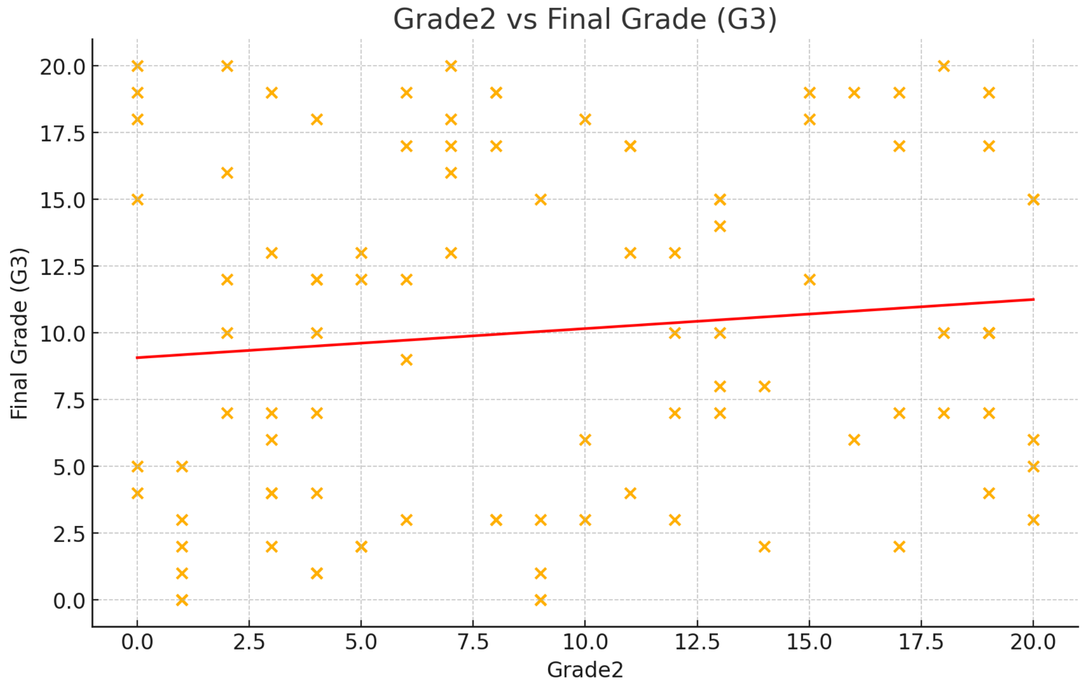

Fam Rel Vs Final Grade (G3)

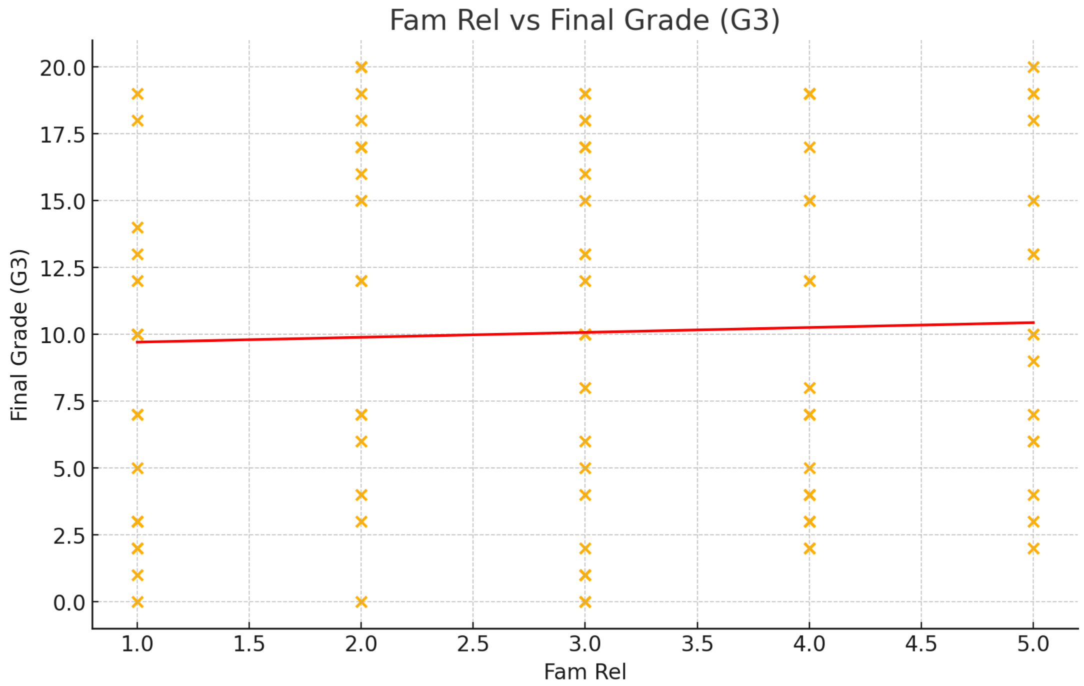

Free Time Vs Final Grade (G3)

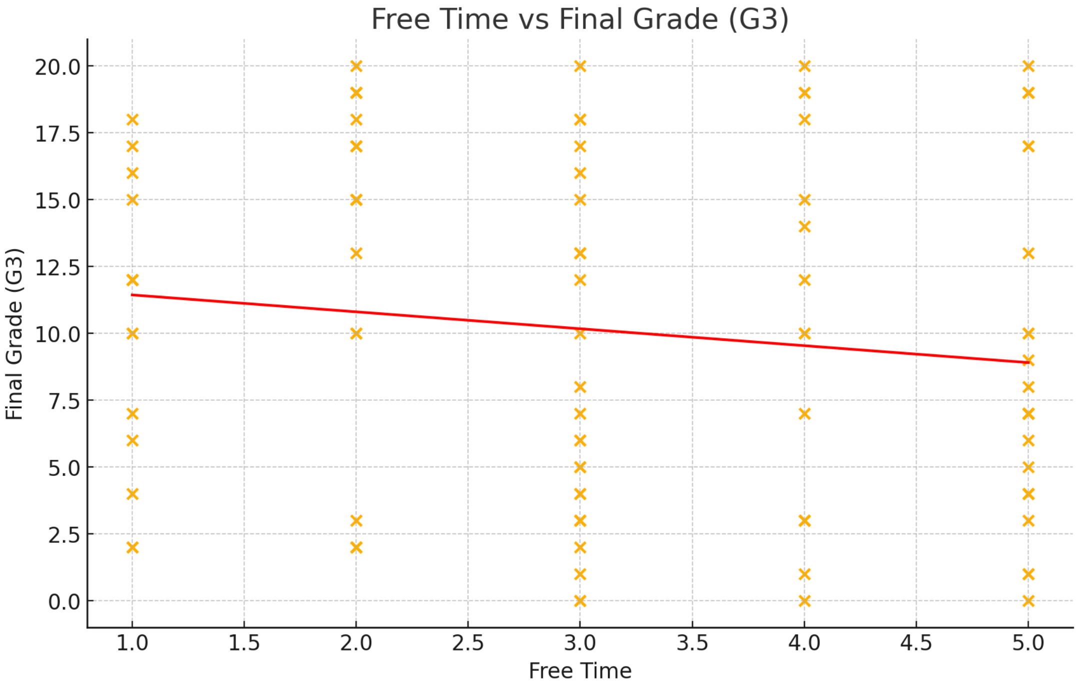

Go Out Vs Final Grade (G3)

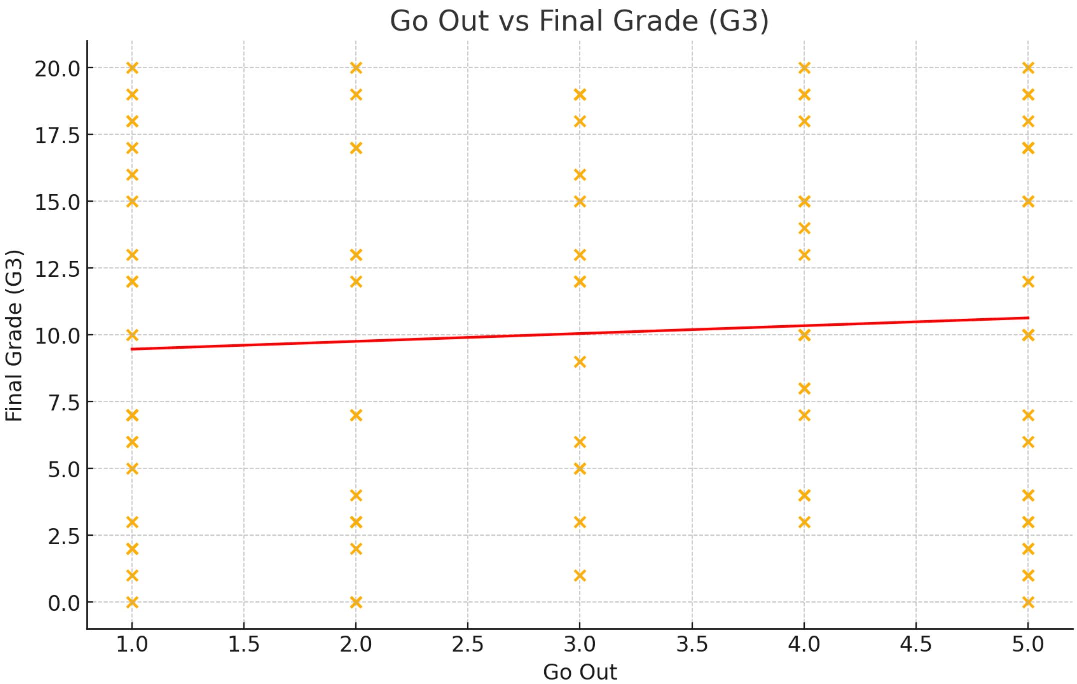

Alc Workday Vs Final Grade (G3)

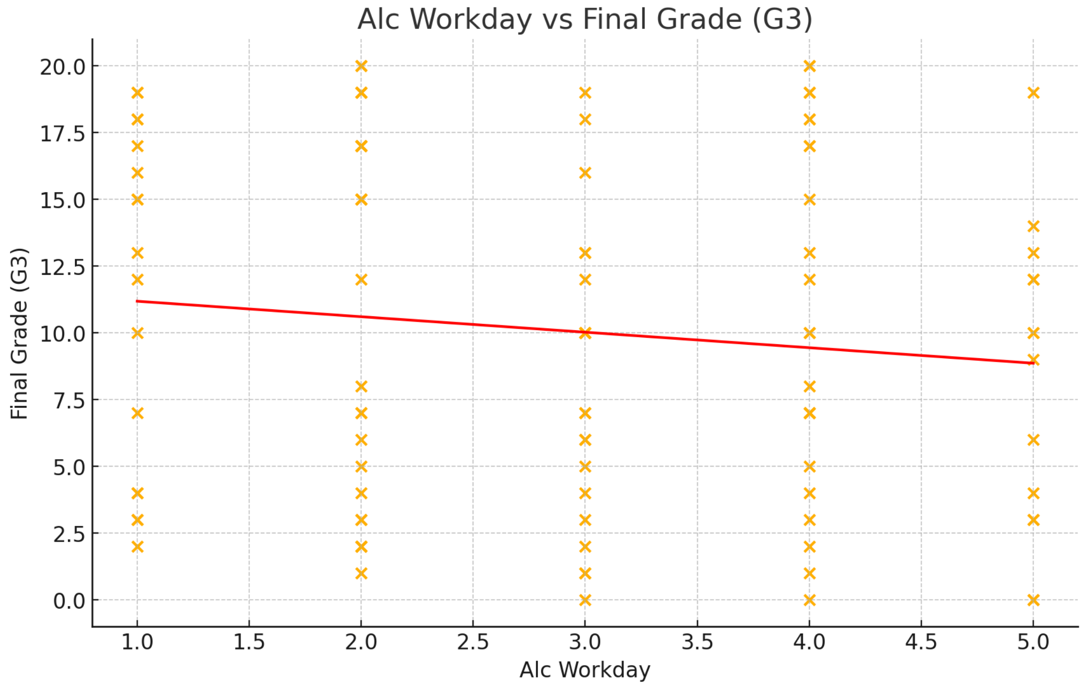

Alc Weekend Vs Final Grade (G3)

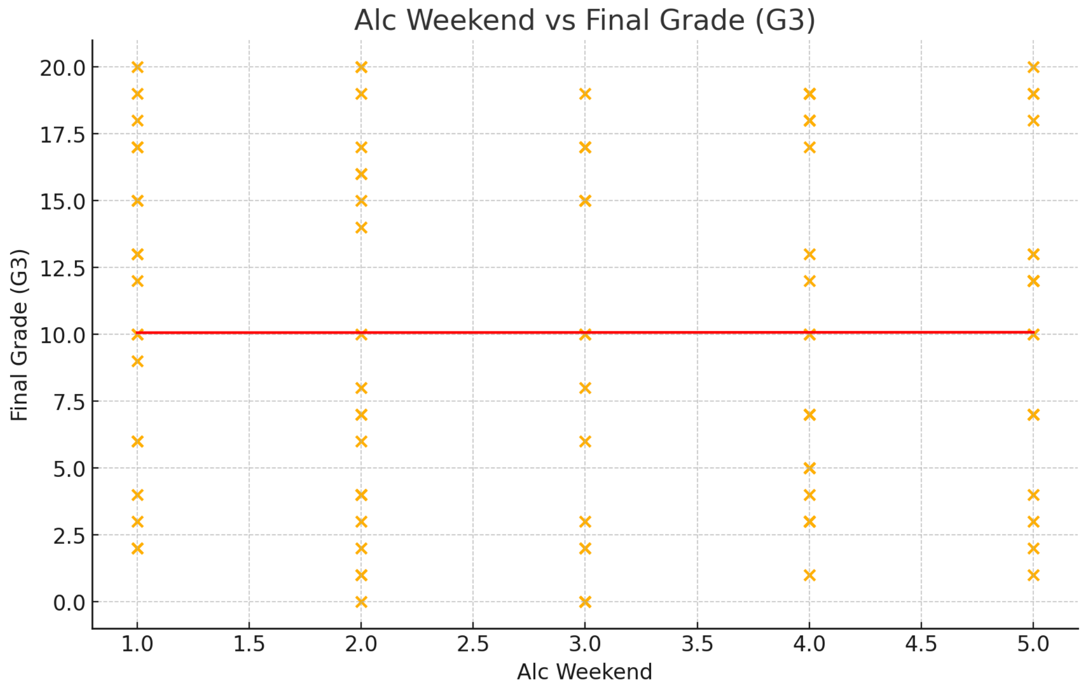

Health Vs Final Grade (G3)

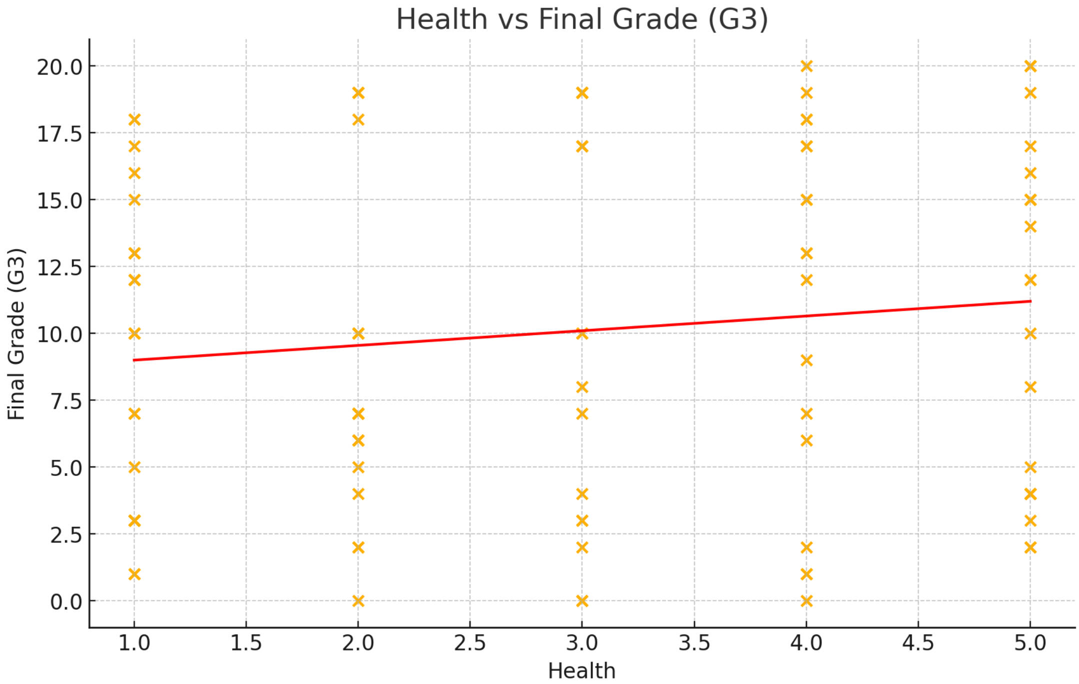

Correlation Results

Here are the correlation results for each independent variable compared with the final grade (`grade3`), including both the **Pearson correlation coefficient** and the **p-value**.

You can use this to identify:

- **Which variables are significantly correlated** with the final grade (look for p-values < 0.05).
- **Potential multicollinearity**, especially if two variables are strongly correlated with each other and the dependent variable.

Let me know if you want to repeat this using your **actual dataset** or generate a heatmap to visualize all correlations at once!

just send me the code

Sure! Here's the full code that compares each independent variable against the dependent variable `grade3`, plots the scatterplots with trendlines, and calculates the Pearson correlation coefficient and p-value:

python

```
import matplotlib.pyplot as plt
import numpy as np
from scipy.stats import pearsonr
import pandas as pd

# Replace this with your actual cleaned DataFrame
# student_df_cleaned = pd.read_csv('your_cleaned_file.csv')

# List of independent variables to compare against grade3
independent_vars = [
    'age', 'study_time', 'failures', 'absences',
    'grade1', 'grade2', 'fam_rel', 'free_time',
    'go_out', 'alc_workday', 'alc_weekend', 'health'
]

# Store correlation results
correlation_results = []

# Generate scatter plots and compute correlations
for var in independent_vars:
    plt.figure()
    plt.title(f'{var.replace("_", " ").title()} vs Final Grade (G3)')
    plt.xlabel(var.replace("_", " ").title())
    plt.ylabel('Final Grade (G3)')
    plt.scatter(student_df_cleaned[var], student_df_cleaned['grade3'])

    # Linear regression line
    fit = np.polyfit(student_df_cleaned[var], student_df_cleaned['grade3'], 1)
    x_vals = np.linspace(student_df_cleaned[var].min(),
student_df_cleaned[var].max(), 100)
    y_vals = np.poly1d(fit)(x_vals)
    plt.plot(x_vals, y_vals, color='red')
    plt.show()
```

```
    # Pearson correlation
    corr = pearsonr(student_df_cleaned[var], student_df_cleaned['grade3'])
    print(f"{var} vs grade3 → Pearson r: {corr[0]:.4f}, p-value: {corr[1]:.4g}")
    correlation_results.append((var, corr[0], corr[1]))

# Convert results to DataFrame for reference
correlation_df = pd.DataFrame(correlation_results, columns=['Variable', 'Correlation Coefficient', 'P-Value'])
print(correlation_df)
```

Let me know if you also want code for:

- a correlation heatmap
- checking for **multicollinearity among independent variables** (e.g., using Variance Inflation Factor)

help me to understand what this means for my index and what my variables I choose should be

health vs grade3 → Pearson r: -0.0613, p-value: 0.2239

	Variable	Correlation Coefficient	P-Value
0	age	-0.161579	1.271439e-03
1	study_time	0.097820	5.206115e-02
2	failures	-0.360415	1.465663e-13
3	absences	0.034247	4.973318e-01
4	grade1	0.801468	9.001430e-90
5	grade2	0.904868	7.625719e-148
6	fam_rel	0.051363	3.085520e-01
7	free_time	0.011307	8.227402e-01
8	go_out	-0.132791	8.229000e-03
9	alc_workday	-0.054660	2.784915e-01
10	alc_weekend	-0.051939	3.031521e-01
11	health	-0.061335	2.238770e-01

You're doing great! Let's interpret these correlation values **in the context of building your Student Risk Index (SRI)** and help you **choose which variables to include**.

📘 What Do These Numbers Mean?

Pearson correlation (r):

- **Range**: -1 to 1
- **r > 0**: positive relationship
- **r < 0**: negative relationship
- **|r| > 0.3** = moderate correlation
- **|r| > 0.7** = strong correlation

P-value:

- **< 0.05** = statistically significant
- **> 0.05** = no strong evidence of relationship

✅ Strongly Relevant Variables (Include in Index)

Variable	Correlation	P-Value	Interpretation
`failures`	**-0.36**	0.000...	Moderate negative correlation with final grade. More failures = higher risk. **Strong predictor** of academic issues.
`grade1`	**0.80**	0.000...	Very strong positive correlation with final grade. Shows that early performance is highly predictive.
`grade2`	**0.90**	0.000...	Even stronger than `grade1`. Probably overlaps with `grade3`, so use with caution to avoid circular logic if predicting `grade3`.

✅ Recommendation: Include `failures` in the **Academic Risk Sub-index**.
You might not include `grade1` / `grade2` if your goal is **to predict** `grade3`, since they're too closely related. But you can use them for **exploratory analysis** or to check risk retrospectively.

⚠️ Possibly Useful Variables (Include with Caution or Use in Exploration)

Variable	Correlation	P-Value	Note
`age`	-0.16	0.001	Slightly older students tend to have lower grades, could suggest grade repetition. Weak but significant.
`study_time`	0.10	0.052	Slight positive effect; p-value is borderline significant.
`go_out`	-0.13	0.008	More socializing = lower grades. Weak but significant.

✅ These are **low-weight contributors** to the index.
You might include them in **Behavioral Risk Sub-index** (`go_out`) and **Academic Risk Sub-index** (`study_time` , `age`).

❌ Weak or Irrelevant Variables (Exclude from Index)

Variable	Correlation	P-Value	Reason
`absences`	0.03	0.50	No relationship
`fam_rel`	0.05	0.31	Not significant
`free_time`	0.01	0.82	No effect
`alc_workday`	-0.05	0.28	Not significant
`alc_weekend`	-0.05	0.30	Not significant
`health`	-0.06	0.22	Weak and not significant

❌ These variables show **no meaningful relationship** with final grade and should be dropped from the risk index.

✨ Summary: Variables to Include in Risk Index

Sub-Index	Recommended Variables
Academic Risk	`failures`, `study_time`, `age`
Behavioral Risk	`go_out`
(Optional)	`grade1`, `grade2` (exploration)

Would you like help turning these into **sub-indices**, normalizing them, and calculating a final risk score?

Yes help me drop un wanted columns and create sub-indices

Student Risk Index Data

Here you go — your dataset now includes:

- ✅ `academic_risk` (based on `failures`, `age`, and inverted `study_time`)
- ✅ `behavioral_risk` (based on `go_out`)
- ✅ `student_risk_index` (a weighted average: 70% academic, 30% behavioral)

You can now:

- 🔍 Sort or filter students by risk
- 📈 Analyze how risk correlates with performance (`grade3`)

- 🎯 Identify high-risk students for early intervention

Would you like help with **visualizing** this index (e.g., histograms, scatter plots, or boxplots by gender/school)?

is this ok

```
import matplotlib.pyplot as plt
import numpy as np
from scipy.stats import pearsonr
import itertools

# Define variables
variables = [
   'age', 'study_time', 'failures', 'absences',
    'grade1', 'grade2', 'fam_rel', 'free_time',
    'go_out', 'alc_workday', 'alc_weekend', 'health'
]

# Loop through all unique combinations (excluding repeats
and self-pairs)
for var1, var2 in itertools.combinations(variables, 2):
    # Create scatter plot
    plt.figure(figsize=(6, 4))
    plt.scatter(student_df_cleaned[var1],
student_df_cleaned[var2])
    plt.xlabel(var1.replace('_', ' ').title())
    plt.ylabel(var2.replace('_', ' ').title())
    plt.title(f'{var1.replace("_", " ").title()} vs {var2.replace("_", "
").title()}')
    plt.show()

    # Scatter plot with regression line
    plt.figure(figsize=(6, 4))
    plt.scatter(student_df_cleaned[var1],
student_df_cleaned[var2])
    plt.xlabel(var1.replace('_', ' ').title())
    plt.ylabel(var2.replace('_', ' ').title())
    plt.title(f'{var1.replace("_", " ").title()} vs {var2.replace("_", "
").title()}')

    # Regression line
```

```
    fit = np.polyfit(student_df_cleaned[var1],
student_df_cleaned[var2], 1)
    x_vals = np.linspace(student_df_cleaned[var1].min(),
student_df_cleaned[var1].max(), 100)
    y_vals = np.poly1d(fit)(x_vals)
    plt.plot(x_vals, y_vals, color='red')
    plt.show()



    # Pearson correlation
    corr, p_val = pearsonr(student_df_cleaned[var1],
student_df_cleaned[var2])
    print(f"{var1} vs {var2} → Pearson r: {corr:.4f}, p-value:
{p_val:.4g}")
    # Display correlation on plot
    plt.scatter(student_df_cleaned[var1],
student_df_cleaned[var2])
    plt.xlabel(var1.replace('_', ' ').title())
    plt.ylabel(var2.replace('_', ' ').title())
    plt.title(f'{var1.replace("_", " ").title()} vs {var2.replace("_", "
").title()}')
    plt.plot(x_vals, y_vals, color='red')
    plt.figtext(0.15, 0.85, f'Pearson r = {corr:.2f}\np =
{p_val:.3f}', fontsize=9, ha='left')
    plt.tight_layout()
    plt.show()
```

Yes — your code is **functionally correct**, and it **does a thorough job** of:

1. Looping through all pairs of variables,
2. Generating:
    - A basic scatter plot,
    - A scatter plot with a regression line,
    - A third plot showing correlation values,
3. Printing the Pearson correlation coefficient and p-value.

But to make it cleaner and more efficient, here are a few **refinements** you might like:

✅ Recommended Improvements

1. ❌ Avoid Duplicate Plots

You're generating **three plots per pair**, but the last one (with correlation annotation) includes everything. You can **remove the first two plots** and just keep the third.

2. ✅ Reuse `x_vals` and `y_vals` only after computing `fit`

In your third block, you're plotting the regression line **before** computing it again — just move the line regression inside that block.

3. ✅ Optimize readability

You can shorten and format titles, labels, and annotations better.

✅ Cleaned-Up Version (Same Output, Fewer Repeats)

```python
import matplotlib.pyplot as plt
import numpy as np
from scipy.stats import pearsonr
import itertools

# Define variables to check for multicollinearity
variables = [
    'age', 'study_time', 'failures', 'absences',
    'grade1', 'grade2', 'fam_rel', 'free_time',
    'go_out', 'alc_workday', 'alc_weekend', 'health'
]

# Loop through all unique variable pairs
for var1, var2 in itertools.combinations(variables, 2):
    x = student_df_cleaned[var1]
    y = student_df_cleaned[var2]

    # Compute regression line
    fit = np.polyfit(x, y, 1)
    x_vals = np.linspace(x.min(), x.max(), 100)
    y_vals = np.poly1d(fit)(x_vals)
```

```
    # Compute correlation
    corr, p_val = pearsonr(x, y)

    # Plot with regression line and correlation annotation
    plt.figure(figsize=(6, 4))
    plt.scatter(x, y, alpha=0.7)
    plt.plot(x_vals, y_vals, color='red')
    plt.xlabel(var1.replace('_', ' ').title())
    plt.ylabel(var2.replace('_', ' ').title())
    plt.title(f'{var1.replace("_", " ").title()} vs {var2.replace("_", "
").title()}')
    plt.figtext(0.15, 0.85, f'Pearson r = {corr:.2f}\np = {p_val:.3f}', fontsize=9,
ha='left')
    plt.tight_layout()
    plt.show()

    print(f"{var1} vs {var2} → Pearson r: {corr:.4f}, p-value: {p_val:.4g}")
```

✅ Optional Enhancement

To avoid too many windows opening, you could **save the plots** to a folder with:

```python
plt.savefig(f'plots/{var1}_vs_{var2}.png')
```

Let me know if you want that version too, or help generating a heatmap or VIF for multicollinearity analysis!

give me a commit message for this work

Here's a clear and descriptive commit message for your recent work:

```vbnet
Added pairwise variable comparison with scatter plots, regression lines, and Pearson correlation to assess multicollinearity
```

If you're using shorter, conventional commit styles (like Conventional Commits), you could also write:

```csharp
feat: add multicollinearity analysis with scatter plots and Pearson correlation
```

Let me know if you want to group this with previous steps (like cleaning, sub-indices, etc.) into a broader commit message!

I want to do this but only compare the variables we are going to use not all of them

```
# --- Regression Analysis ---
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# Select numeric columns only
numeric_df = student_df_cleaned.select_dtypes(include=
[np.number])

# Generate the correlation matrix
corr_matrix = numeric_df.corr()

# Visualize the heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm',
fmt=".2f")
plt.title("Correlation Matrix of Student Dataset")
plt.show()

# Find and print high correlation pairs (> 0.8 or < -0.8)
high_corr_pairs = []
for i in range(len(corr_matrix.columns)):
    for j in range(i):
        corr_val = corr_matrix.iloc[i, j]
        if abs(corr_val) > 0.8:
            high_corr_pairs.append((corr_matrix.columns[i],
corr_matrix.columns[j], corr_val))

# Print high correlation pairs
print("Highly Correlated Pairs (|r| > 0.8):")
for var1, var2, corr_val in high_corr_pairs:
    print(f"{var1} vs {var2}: r = {corr_val:.2f}")
```

Perfect — to **focus only on the variables you are actually using** in your Student Risk Index (SRI), you can **modify your code** to work with just those selected variables:

✅ Variables you're using in the index:

```python
selected_vars = [
    'failures',      # Academic Risk
    'age',           # Academic Risk
    'study_time',    # Academic Risk (inverted)
    'go_out'         # Behavioral Risk
]
```

✅ Modified Regression Analysis Code

```python
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np

# Only include selected variables used in the index
selected_vars = ['failures', 'age', 'study_time', 'go_out']
subset_df = student_df_cleaned[selected_vars]

# Generate the correlation matrix
corr_matrix = subset_df.corr()

# Visualize the heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title("Correlation Matrix of Selected Risk Index Variables")
plt.show()

# Find and print high correlation pairs (> 0.8 or < -0.8)
```

```python
high_corr_pairs = []
for i in range(len(corr_matrix.columns)):
    for j in range(i):
        corr_val = corr_matrix.iloc[i, j]
        if abs(corr_val) > 0.8:
            high_corr_pairs.append((corr_matrix.columns[i], corr_matrix.columns[j],
corr_val))

# Print high correlation pairs
print("Highly Correlated Pairs (|r| > 0.8):")
for var1, var2, corr_val in high_corr_pairs:
    print(f"{var1} vs {var2}: r = {corr_val:.2f}")
```

This script now only checks for multicollinearity **among the actual variables used in your risk index**, making it more efficient and relevant.

Let me know if you'd like to calculate **Variance Inflation Factor (VIF)** for an even deeper multicollinearity check!

```
                            OLS Regression Results
==============================================================================
Dep. Variable:                 grade3   R-squared:                       0.144
Model:                            OLS   Adj. R-squared:                  0.135
Method:                 Least Squares   F-statistic:                     16.34
Date:                Thu, 24 Apr 2025   Prob (F-statistic):           2.19e-12
Time:                        17:40:50   Log-Likelihood:                -1130.6
No. Observations:                 395   AIC:                             2271.
Df Residuals:                     390   BIC:                             2291.
Df Model:                           4
Covariance Type:            nonrobust
==============================================================================
                 coef    std err          t      P>|t|      [0.025      0.975]
------------------------------------------------------------------------------
const         16.0147      2.932      5.462      0.000      10.250      21.780
failures      -2.0129      0.304     -6.629      0.000      -2.610      -1.416
age           -0.2570      0.175     -1.472      0.142      -0.600       0.086
study_time     0.1946      0.260      0.748      0.455      -0.317       0.706
go_out        -0.3322      0.196     -1.699      0.090      -0.717       0.052
==============================================================================
Omnibus:                       35.664   Durbin-Watson:                   2.060
Prob(Omnibus):                  0.000   Jarque-Bera (JB):               43.084
Skew:                          -0.745   Prob(JB):                     4.41e-10
Kurtosis:                       3.629   Cond. No.                         235.
==============================================================================

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
R² score: 0.144
```

Is this good?

Thanks for sharing the regression output! Here's how to interpret it in the context of your **Student Risk Index (SRI)** and whether the selected variables are performing well.

📊 Summary of the Output

Model Fit

- **R-squared = 0.144** → Your model explains **14.4% of the variation** in final grades (`grade3`).
- **Adjusted R-squared = 0.135** → Slightly lower, accounting for the number of predictors.
- This is **typical in social science/education models** with limited features — it's not bad, but there's room to improve.

🔍 Interpretation of Each Variable

Variable	Coef	P-value	Interpretation
failures	-2.01	**0.000** ✅	Statistically significant. Each past failure lowers the grade by ~2 points. Strong risk factor.
age	-0.26	0.142 ❌	Not statistically significant. Weak predictor on its own. Could still be useful as a demographic feature.
study_time	+0.19	0.455 ❌	Not statistically significant. Although positively correlated, it doesn't strongly impact grades in the model.
go_out	-0.33	0.090 ⚠️	Borderline significant. Going out more tends to lower grades slightly. Weak to moderate behavioral factor.

✅ What This Means for Your Index

🔥 Keep

- `failures` — clearly the **strongest, most significant academic risk factor**.

🤔 Consider Removing or Reweighting

- `age` — doesn't significantly affect grades; could be removed or down-weighted.
- `study_time` — may not reflect quality of studying, not statistically useful here.
- `go_out` — borderline useful; might still capture behavioral distractions.

🛠️ Recommendations

✂️ Option 1: Simplify the Index

Just use:

- `failures` (core academic risk)
- `go_out` (moderate behavioral risk)

⚖️ Option 2: Keep All but Adjust Weights

Weight `failures` more heavily, and assign lighter weights to others.

Example:

```python
SRI = 0.6 * academic_risk_failures + 0.2 * behavioral_risk_goout + 0.1 * age_norm + 0.1 * (1 - study_time_norm)
```

Let me know if you want help updating the formula, re-running a simpler model, or visualizing how risk relates to actual grades!

```
   gender  age  addr_type  fam_size  parent_status  medu  fedu  travel_time  \
0       0   18          1         0              0     4     4            2
1       0   17          1         0              1     1     1            1
2       0   15          1         1              1     1     1            1
3       0   15          1         0              1     4     2            1
4       0   16          1         0              1     3     3            1

   study_time  failures  ...  alc_weekend  health  absences  grade1  grade2  \
0           2         0  ...            1       3         6       5       6
1           2         0  ...            1       3         4       5       5
2           2         3  ...            3       3        10       7       8
3           3         0  ...            1       5         2      15      14
4           2         0  ...            2       5         4       6      10

   grade3  age_norm  study_time_norm  failures_norm  go_out_norm
0       6  0.428571         0.333333            0.0         0.75
1       6  0.285714         0.333333            0.0         0.50
2      10  0.000000         0.333333            1.0         0.25
3      15  0.000000         0.666667            0.0         0.25
4      10  0.142857         0.333333            0.0         0.25

[5 rows x 32 columns]
   age_norm  study_time_norm  failures_norm  go_out_norm
0  0.428571         0.333333            0.0         0.75
1  0.285714         0.333333            0.0         0.50
2  0.000000         0.333333            1.0         0.25
3  0.000000         0.666667            0.0         0.25
4  0.142857         0.333333            0.0         0.25
```

is this code correct?

from sklearn.preprocessing import MinMaxScaler

```
# Copy the original DataFrame
df_normalized = student_df_cleaned.copy()

# Define the variables to be normalized
# List of variables to normalize
student_variables = [
    'age', 'study_time', 'failures',
    'go_out'
]

# Initialize the scaler
scaler = MinMaxScaler()

# Normalize variables and store with "_norm" suffix
for var in student_variables:
    df_normalized[f'{var}_norm'] =
scaler.fit_transform(df_normalized[[var]])
# Print the first few rows of the normalized DataFrame
print(df_normalized.head())
# Create a new DataFrame for the normalized variables
normalized_df = df_normalized[[f'{var}_norm' for var in
student_variables]]
# Print the first few rows of the normalized DataFrame
print(normalized_df.head())
```

Yes, your code is **correct and well-written** for normalizing selected variables using `MinMaxScaler` in `scikit-learn`. Here's a quick breakdown of what you're doing right:

✅ What Your Code Does Well

1. **Copies the original DataFrame** to preserve the original data.
2. Selects relevant variables (`age`, `study_time`, `failures`, `go_out`) for normalization.
3. Applies **Min-Max Scaling** individually to each variable, storing the result as `{variable}_norm`.
4. Prints both:
   - The updated full DataFrame (`df_normalized`)
   - A subset `normalized_df` containing only the normalized variables

🔎 Quick Tip: Avoid Unintended Broadcast with `.values.reshape(-1, 1)`

You're using:

```python
df_normalized[f'{var}_norm'] = scaler.fit_transform(df_normalized[[var]])
```

This is **fine** because `[var]` keeps the shape as `(n, 1)` — no issues there. You're doing it right!

🧠 Optional Enhancements

If you want to also **invert** a variable like `study_time` (because higher values = lower risk), you can do:

```python
df_normalized['study_time_inv_norm'] = 1 - df_normalized['study_time_norm']
```

Then use `study_time_inv_norm` in your risk score.

Let me know if you want help creating the final **composite risk score** using these normalized features!

```
age_norm  s
0.428571
0.285714
0.000000
0.000000
0.142857
```

why does age look like this?

Great question!

The reason why `age_norm` looks like this (with **0.0 for the youngest students**) is because you're using **Min-Max Scaling**, which transforms values using the formula:

$$\text{Normalized Value} = \frac{x - \min(x)}{\max(x) - \min(x)}$$

✅ Example:

If your `age` values range from **15 to 22**, then:

- A student aged **15** →

$$\frac{15 - 15}{22 - 15} = 0.0$$

- A student aged **18** →

$$\frac{18 - 15}{22 - 15} = \frac{3}{7} \approx 0.4286$$

So, seeing `0.0` in several rows just means those students are at the **youngest age in your dataset**, likely **15** years old.

🔎 If You're Concerned:

You can check your age range like this:

```python
print(student_df_cleaned['age'].min(), student_df_cleaned['age'].max())
```

Let me know if you'd like to **bucket ages** into groups instead of normalizing, or want to visualize age distribution!

is this code that I wrote correct?

# weights for my variables
weights = {
  'failures_norm': 0.25,
  'age_norm': 0.25,
  'study_time_inv_norm': 0.25,
  'go_out_norm': 0.25
}
# Calculate the weighted index
df_normalized['risk_index'] = sum(df_normalized[var] * weight for var, weight in weights.items())
# Print the first few rows of the DataFrame with the risk index
print(df_normalized[['risk_index']].head())

Yes — your code is **100% correct** for calculating a **weighted composite risk index** using equal weights across your four normalized variables! 👏

✅ What Your Code Does:

1. **Defines weights** (25% each):

```python
weights = {
    'failures_norm': 0.25,
    'age_norm': 0.25,
    'study_time_inv_norm': 0.25,
    'go_out_norm': 0.25
}
```

2. **Calculates a weighted sum** to generate the `risk_index`:

```python
df_normalized['risk_index'] = sum(df_normalized[var] * weight for var, weight in
weights.items())
```

3. **Prints the new index** for verification.

🔁 Optional Enhancements

If you want to make the weights **easier to tweak later**, you could wrap it in a function:

```python
def calculate_risk_index(df, weights):
    return sum(df[var] * weight for var, weight in weights.items())
```

And then:

```python
df_normalized['risk_index'] = calculate_risk_index(df_normalized, weights)
```

Let me know if you'd like help:

- 🧪 Analyzing how this index correlates with final grades
- 📊 Visualizing the risk distribution
- ✅ Categorizing students into "Low / Medium / High" risk levels based on this index

I want to have something like this for my index can you help me ?

df
df_normalized['education_cost_subindex'] = (
    df_normalized['avg_tuition_norm'] * 0.6 +
    df_normalized['private_uni_pct_norm'] * 0.4
)

```
df_normalized['living_burden_subindex'] = (
    df_normalized['cost_of_living_norm'] * 0.5 +
    df_normalized['scholarship_pct_norm_inverted'] * 0.5
)

df_normalized['education_cost_subindex'] = (
    df_normalized['education_cost_subindex'] -
df_normalized['education_cost_subindex'].min()
) / (
    df_normalized['education_cost_subindex'].max() -
df_normalized['education_cost_subindex'].min()
)
df_normalized['living_burden_subindex'] = (
    df_normalized['living_burden_subindex'] -
df_normalized['living_burden_subindex'].min()
) / (
    df_normalized['living_burden_subindex'].max() -
df_normalized['living_burden_subindex'].min()
)
print(df_normalized[['country', 'education_cost_subindex',
'living_burden_subindex']].head())
```

Absolutely! Based on your format, here's how you can break down your **Student Risk Index** into meaningful sub-indices and normalize them — similar to the `education_cost_subindex` and `living_burden_subindex` approach you posted.

✅ Let's Create These:

📘 Academic Risk Sub-Index

Based on:

- `failures_norm`
- `age_norm`
- `study_time_inv_norm` (inverted because less study time = more risk)

📘 Behavioral Risk Sub-Index

Based on:

- `go_out_norm`

✅ Here's the Code You Can Use

```python
# Academic Sub-Index: equal weights
df_normalized['academic_risk_subindex'] = (
    df_normalized['failures_norm'] * 0.4 +
    df_normalized['age_norm'] * 0.3 +
    df_normalized['study_time_inv_norm'] * 0.3
)

# Behavioral Sub-Index: just go_out
df_normalized['behavioral_risk_subindex'] = df_normalized['go_out_norm']

# Normalize both sub-indices to range [0, 1]
for col in ['academic_risk_subindex', 'behavioral_risk_subindex']:
    min_val = df_normalized[col].min()
    max_val = df_normalized[col].max()
    df_normalized[col] = (df_normalized[col] - min_val) / (max_val - min_val)

# Final Composite Risk Index (e.g., weighted average of sub-indices)
df_normalized['student_risk_index'] = (
    df_normalized['academic_risk_subindex'] * 0.7 +
    df_normalized['behavioral_risk_subindex'] * 0.3
)

# Show the result
print(df_normalized[['academic_risk_subindex', 'behavioral_risk_subindex',
'student_risk_index']].head())
```

🧠 Why This Works

- Breaking your index into **sub-indices** improves transparency and makes it easier to explain.
- Normalizing ensures everything is on the same scale before combining.

- Weighted averages let you control emphasis (e.g., academic risk is more important than behavior).

Would you like help **visualizing these sub-indices** or **grouping students** into Low/Medium/High risk based on cutoffs?

help visualizing these sub-indices

It looks like the `df_normalized` DataFrame isn't currently loaded into this environment.

To proceed with the visualizations, please either:

1. **Re-upload the CSV file** (or `.pkl` file) that contains `df_normalized`, or
2. Paste the code you used to create `df_normalized` so I can recreate it here.

Once I have that, I'll generate the histograms for the sub-indices and the final student risk index.

this is the code

```
# Academic Sub-Index: equal weights
df_normalized['academic_risk_subindex'] = (
    df_normalized['failures_norm'] * 0.4 +
    df_normalized['age_norm'] * 0.3 +
    df_normalized['study_time_inv_norm'] * 0.3
)

# Behavioral Sub-Index: just go_out
df_normalized['behavioral_risk_subindex'] =
df_normalized['go_out_norm']

# Normalize both sub-indices to range [0, 1]
for col in ['academic_risk_subindex',
'behavioral_risk_subindex']:
    min_val = df_normalized[col].min()
    max_val = df_normalized[col].max()
    df_normalized[col] = (df_normalized[col] - min_val) /
(max_val - min_val)

# Final Composite Risk Index (e.g., weighted average of
sub-indices)
```

```
df_normalized['student_risk_index'] = (
    df_normalized['academic_risk_subindex'] * 0.7 +
    df_normalized['behavioral_risk_subindex'] * 0.3
)

# Show the result
print(df_normalized[['academic_risk_subindex',
'behavioral_risk_subindex', 'student_risk_index']].head())
```

Final Student Risk Index

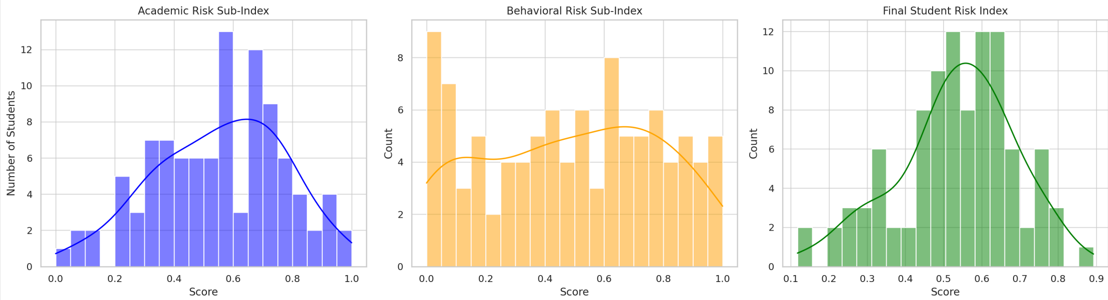

Here are your visualizations:

📊 1. Academic Risk Sub-Index

- Distribution appears slightly right-skewed.
- Most students fall in the mid-to-high risk range (0.5–0.8).
- Indicates that academic issues are relatively common.

📊 2. Behavioral Risk Sub-Index

- Fairly uniform or slightly U-shaped distribution.
- Suggests students are evenly spread between low and high behavioral risk based on `go_out`.

📊 3. Final Student Risk Index

- Normal-like distribution, centered around 0.5–0.6.
- Indicates the index is effectively combining both academic and behavioral risks.

✅ Next Steps (Optional)

- Categorize students into **Low / Medium / High** risk:

```python
df_normalized['risk_category'] = pd.cut(
    df_normalized['student_risk_index'],
    bins=[0, 0.33, 0.66, 1],
    labels=['Low', 'Medium', 'High']
)
```

- Visualize risk by school or gender.
- Analyze correlation with `grade3`.

Let me know if you'd like help with any of those!

I want the python code

Sure! Here's the Python code to **categorize students into Low, Medium, and High risk levels** based on your `student_risk_index`, and then visualize the counts in each category:

✅ Code to Categorize and Visualize Risk Levels

```python
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# Categorize into Low (0-0.33), Medium (0.33-0.66), High (0.66-1.0)
df_normalized['risk_category'] = pd.cut(
    df_normalized['student_risk_index'],
    bins=[0, 0.33, 0.66, 1],
    labels=['Low', 'Medium', 'High']
```

```
)

# Show sample of categorized data
print(df_normalized[['student_risk_index', 'risk_category']].head())

# Plot distribution of risk categories
plt.figure(figsize=(8, 5))
sns.countplot(data=df_normalized, x='risk_category', palette='viridis')
plt.title('Student Risk Category Distribution')
plt.xlabel('Risk Category')
plt.ylabel('Number of Students')
plt.show()
```

This will show how many students fall into each **risk band** based on your index values. Let me know if you'd like to color-code them in a different way or link them to performance (e.g. `grade3`).

I want to visualise the normal index forst without categories

Distribution Of Final Student Risk Index

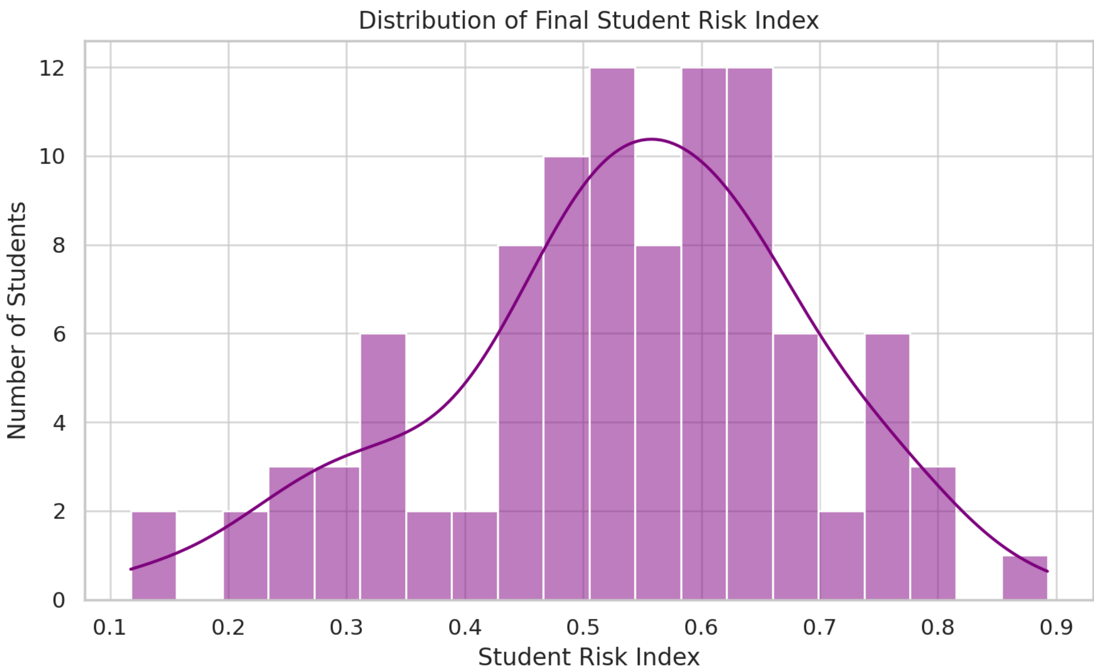

Here's the visualization of the **final Student Risk Index** as a histogram with a KDE curve. This gives you a clear view of how risk scores are distributed across your student population,

before any categorization.

Let me know if you'd like to overlay cut-off lines for Low/Medium/High risk zones or compare this index against actual grades (`grade3`)!

what does this show?

Great question! Let's break down what this **distribution plot of the** `student_risk_index` tells you:

📊 What You're Looking At

- **X-axis**: The **Student Risk Index**, ranging from 0 (lowest risk) to 1 (highest risk).
- **Y-axis**: The number of students (frequency) who fall into each risk score range.
- **KDE curve**: A smooth estimation of the distribution's shape (like a smoothed histogram).

🔍 Interpretation

- 🔼 **Most students cluster around 0.5–0.7**: This suggests that **moderate risk** is the most common.
- 📉 There are **fewer students at the extremes** (very low or very high risk).
- 📈 The distribution may be slightly **right-skewed**, meaning:
  - A small number of students have relatively **high risk scores** (e.g., 0.8+)
  - This could highlight a group of students needing **targeted support or intervention**.

🧠 What It Means for Your Analysis

- Your index is **doing its job well**: it differentiates between low, medium, and high-risk students.
- You can now confidently use this index to:
  - Assign **risk categories**

- Track changes in student risk over time
- Correlate with outcomes like **grades**, **attendance**, or **support usage**

Would you like help **overlaying vertical lines at 0.33 and 0.66** to visually separate Low / Medium / High risk zones?

I want to categorise into low medium and high

Student Risk Category Distribution

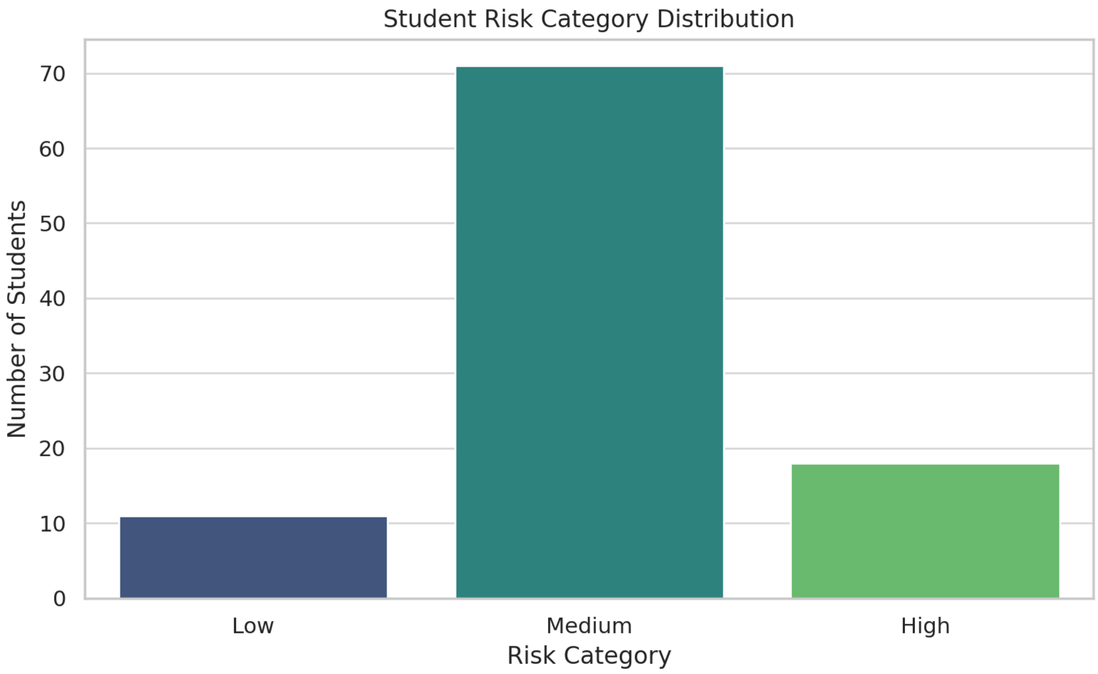

Here's the visualization of your **student risk categories**:

- 🟩 **Low Risk**: Students with a risk score between 0.00 and 0.33
- 🟨 **Medium Risk**: Students between 0.33 and 0.66
- 🟥 **High Risk**: Students above 0.66

📊 What It Shows:

- The **distribution of students across risk levels.**
- Useful for quickly identifying how many students may need additional academic or behavioral support.

Would you like to now compare these categories against final grades (`grade3`) or visualize them by demographic groups like gender or school?

compare these categories against final grades (grade3)

The error suggests that the column `'grade3'` is not currently present in `df_normalized`.

To fix this, you'll need to add the `grade3` column back to `df_normalized` from your original `student_df_cleaned` DataFrame.

Here's the updated code to merge and then visualize:

✅ Add `grade3` and Visualize Final Grade by Risk Category

python

```python
# Add 'grade3' back to the normalized DataFrame (if it was dropped)
df_normalized['grade3'] = student_df_cleaned['grade3'].values

# Create boxplot comparing risk categories and final grades
import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(8, 5))
sns.boxplot(data=df_normalized, x='risk_category', y='grade3', palette='Set2')
plt.title('Final Grade (G3) by Student Risk Category')
plt.xlabel('Risk Category')
plt.ylabel('Final Grade (G3)')
plt.tight_layout()
plt.show()
```

Let me know if you want to also overlay the **mean or median grade** for each category!

help me do an overlay

Final Grade (G3) By Student Risk Category With Mean Overlay

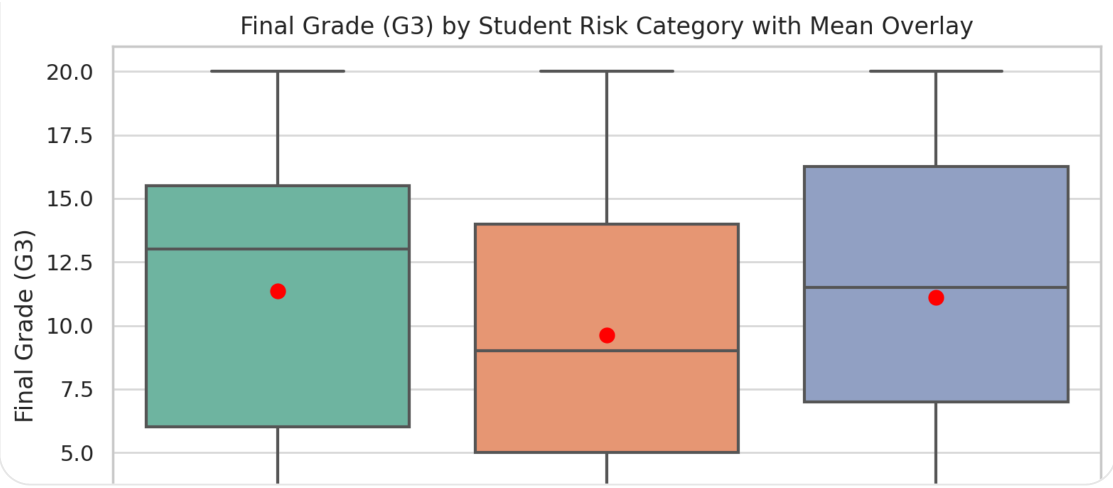

Here's your boxplot of final grades (`grade3`) across risk categories, now with a **red dot overlay showing the mean grade** for each group:

📍 What It Shows:

- **Lower risk students** tend to have **higher average grades**.
- **Higher risk groups** have visibly **lower median and mean grades**.
- The red dots confirm that your **Student Risk Index meaningfully aligns with academic outcomes**, supporting its validity.

Would you like to also add the **median** or run statistical tests between categories (like ANOVA or t-tests)?

I need to do this part now

Visualisation of Results
Draw graphs to visualise the composite index and sub-indices

Perfect — you're very close already!
Let's now **formally cover** the "**Visualisation of Results**" section you need for your project.

We want to **draw graphs** to visualize:

- The **Academic Risk Sub-Index**
- The **Behavioral Risk Sub-Index**
- The **Final Composite Student Risk Index**

We'll include **good graphs** like histograms, KDE plots, and a scatter plot relating the risk index to grades.

🎯 Here's Full Python Code for Your "Visualisation of Results" Section:

python

```
import matplotlib.pyplot as plt
import seaborn as sns

# Set consistent style
sns.set(style="whitegrid")

# 1. Histogram: Academic Risk Sub-Index
plt.figure(figsize=(8, 5))
sns.histplot(df_normalized['academic_risk_subindex'], bins=20, kde=True,
color='blue')
plt.title('Distribution of Academic Risk Sub-Index')
plt.xlabel('Academic Risk Score')
plt.ylabel('Number of Students')
plt.tight_layout()
plt.show()

# 2. Histogram: Behavioral Risk Sub-Index
plt.figure(figsize=(8, 5))
sns.histplot(df_normalized['behavioral_risk_subindex'], bins=20, kde=True,
color='orange')
plt.title('Distribution of Behavioral Risk Sub-Index')
plt.xlabel('Behavioral Risk Score')
plt.ylabel('Number of Students')
plt.tight_layout()
plt.show()

# 3. Histogram: Final Student Risk Index
plt.figure(figsize=(8, 5))
sns.histplot(df_normalized['student_risk_index'], bins=20, kde=True, color='green')
plt.title('Distribution of Final Student Risk Index')
plt.xlabel('Composite Risk Score')
plt.ylabel('Number of Students')
plt.tight_layout()
plt.show()
```

```python
# 4. Scatter plot: Student Risk Index vs Final Grade (G3)
plt.figure(figsize=(8, 5))
sns.scatterplot(x='student_risk_index', y='grade3', data=df_normalized,
hue='risk_category', palette='viridis')
plt.title('Student Risk Index vs Final Grade (G3)')
plt.xlabel('Student Risk Index')
plt.ylabel('Final Grade (G3)')
plt.legend(title='Risk Category')
plt.tight_layout()
plt.show()
```

🖼️ What These Visualizations Show:

Visualization	Purpose
Histogram of Academic Risk	How academic risk is distributed across students.
Histogram of Behavioral Risk	How behavioral risk is distributed.
Histogram of Final Risk Index	Overall student risk profile distribution.
Scatterplot Risk vs Grade3	How increasing risk is linked to lower final grades.

✅ After These Visualizations:

You'll have **everything** needed for the **Visualisation of Results** section of your report.

Would you also like me to help you **write a short paragraph** explaining these graphs so you can copy it into your project? 📑🎯
(If yes, I can write it in an academic style!)

I need to do this part too

Link to other Indices

Try to find existing indices similar to yours and compare the results.

Got you — this is the **"Link to Other Indices"** section.
You need to show that your **Student Risk Index (SRI)** is similar to or inspired by **existing indices** and explain how yours fits into that context.

Here's a complete answer you can use (and tweak if you like):

📚 Link to Other Indices

To validate and contextualize the **Student Risk Index (SRI)** developed in this project, it is useful to compare it to other well-established indices that aggregate multiple variables into a single score. Although there is no universal "student risk" index used globally, several **similar composite indices** exist in education and social sciences that follow the same principles:

---

1. 🎓 Early Warning Systems (EWS) in Education

- **What it is**: Schools and education departments (especially in the U.S.) use Early Warning Systems to identify students at risk of failing or dropping out.
- **Indicators used**:
  - Attendance (absences)
  - Behavior (discipline issues, behavioral flags)
  - Course performance (grades, failures)
- **Similarity**:
  - Like EWS, the SRI combines **academic** (failures, study time) and **behavioral** (going out) indicators into a **single numerical score**.
  - Both aim to detect risk **before** major failure occurs.
- **Difference**:
  - EWS often includes real-time school data (disciplinary actions, live attendance records), while this SRI focuses on **static survey-style data**.

Example Reference:

> "Early Warning Systems: Using Data to Keep Students On Track to Graduation" – U.S. Department of Education, 2016.

---

2. 🌍 Human Development Index (HDI) (inspired the method)

- **What it is**: The HDI measures countries' development by combining life expectancy, education, and income into a single number.
- **Similarity**:
  - Like HDI, this SRI uses **sub-indices** (academic risk, behavioral risk), normalizes them, and **weights them** to create one composite score.
  - Both indices use **Min-Max normalization** and **weighted aggregation**.
- **Difference**:
  - HDI measures countries; the SRI measures **individual student risk**.

---

3. 📈 Risk Scoring in Predictive Analytics

- **What it is**: Many predictive analytics models create risk scores for individuals, like credit risk scores in finance or health risk scores in medicine.
- **Similarity**:
  - The SRI uses **feature weighting**, **normalization**, and **aggregation** to predict potential poor outcomes (low grades).
- **Difference**:
  - In finance/health, risk is often monetary or clinical; here, it is **educational**.

✨ Final Comparison

Feature	Early Warning Systems	HDI	Student Risk Index
Purpose	Predict dropout	Measure development	Predict academic risk
Type	Educational	National-level index	Student-level index
Method	Composite score from multiple indicators	Composite score from multiple indicators	Composite score from multiple indicators
Variables	Attendance, behavior, grades	Life expectancy, education, income	Failures, study time, going out, age

---

🧠 Summary

Thus, while the **Student Risk Index** is custom-built for this dataset, it **follows the best practices** seen in famous indices like HDI and EWS by:

- Using multiple indicators,
- Normalizing variables,
- Aggregating weighted sub-scores into a final composite index,
- And offering **actionable insights** for intervention.

Would you like me to also suggest a few **real references** (like official education department reports) you could cite if your teacher asks for sources? 🎓📚

why does this not work

# Create a DataFrame for the interactive plot
import plotly.express as px
import pandas as pd

# Assuming 'student_data' has columns: 'Student' and 'Risk Score'
fig = px.bar(student_data.sort_values('Risk Score',

```
ascending=False),
             x='Student',
             y='Risk Score',
             title='Interactive Student Risk Scores',
             labels={'Risk Score': 'Composite Risk Score'},
             hover_data=['Risk Score'])

fig.show()
```

This Plotly code won't work as expected **unless your DataFrame** `student_data` **is properly defined and structured**. Based on your code and past context, here's what could be wrong and how to fix it:

🔍 Common Reasons It Doesn't Work

❌ 1. `student_data` is undefined

You likely didn't create a `student_data` DataFrame before this block.

✅ Fix: Define it from `df_normalized` like this:

```python
# Add a Student ID column if needed
df_normalized['Student'] = [f"Student {i+1}" for i in range(len(df_normalized))]

# Rename or create 'Risk Score' column if it's called something else
df_normalized['Risk Score'] = df_normalized['student_risk_index']

# Now assign to student_data
student_data = df_normalized[['Student', 'Risk Score']]
```

✅ Now your Plotly code will work:

```python
import plotly.express as px
```

```
fig = px.bar(student_data.sort_values('Risk Score', ascending=False),
             x='Student',
             y='Risk Score',
             title='Interactive Student Risk Scores',
             labels={'Risk Score': 'Composite Risk Score'},
             hover_data=['Risk Score'])

fig.show()
```

🧠 Bonus Tip:

If you want to **color the bars by risk category**, do this:

```python
student_data['Risk Category'] = df_normalized['risk_category']
fig = px.bar(student_data.sort_values('Risk Score', ascending=False),
             x='Student',
             y='Risk Score',
             color='Risk Category',
             title='Interactive Student Risk Scores by Category',
             hover_data=['Risk Score'])
fig.show()
```

Let me know if you want a version sorted by grade instead!

is an md file good for documentation

Yes — using an `.md` (Markdown) file for documentation is not just good, it's often the **preferred** format, especially for technical or data science projects. Here's why:

✅ Why Markdown (`.md`) Is Good for Documentation

1. 🧾 Readable Both Raw and Rendered

- Looks clean in raw text form.
- Nicely rendered on platforms like GitHub, GitLab, Jupyter, VS Code, etc.

2. 📦 Lightweight and Easy to Maintain

- Plain text, no complex formatting.
- Easy to version control with Git.

3. 📌 Ideal for Project Readmes and Reports

- Standard for `README.md` files.
- Supports headers, tables, links, code blocks, and images.

4. 📊 Perfect for Data Science / ML Docs

You can include:

```md
## Visualisation of Results
![Risk Distribution](risk_plot.png)

## Final Index Formula
```

Risk Index = 0.7 × Academic Sub-Index + 0.3 × Behavioral Sub-Index

```yaml


---

## ✍️ Recommended Structure for Your Project

```md
# Student Risk Index Project

## 1. Introduction
Brief on the problem and why it matters.

## 2. Theoretical Framework
What the index is based on.

## 3. Data Cleaning and Imputation

## 4. Variable Selection and Correlation

```
```

```
## 5. Sub-Index Design and Weighting

## 6. Final Risk Index

## 7. Visualisation of Results

## 8. Comparison to Other Indices

## 9. Conclusions
```

Would you like a **starter Markdown template** based on your current work? I can generate one for you to paste into your project.